

MUHAMMAD HUSNAIN MUBARIK

Website: <https://husnainmubarik.github.io/>

Address: 904 W Stoughton Ave. APT 11, Urbana, IL, 61801

Email: mubarik3@illinois.edu, Phone: +1-217-778-8784

RESEARCH INTERESTS

Computer Architecture, Hardware Acceleration of Machine Learning Applications, Hardware for Graphics (Conventional, Neural and Generative), Hardware Acceleration of Generative AI, Printed Computer Systems, Emerging Technologies and Applications.

EDUCATION

University of Illinois at Urbana Champaign (UIUC), USA. *Jan 2019 - Present*
PhD - Electrical and Computer Engineering (ECE)
Advisor: [Prof. Rakesh Kumar](#)

National University of Sciences and Technology (NUST), Pakistan. *Sep 2013 - July 2017*
Bachelors of Science in Electrical Engineering.
CGPA: 3.95/4.00, **summa cum laude**.

PUBLICATION SUMMARY AND IMPACT

- 5× (ISCA, MICRO), 1× (DATE)
- 1× ISCA Retrospective (1996-2020) Selection
- 1× IEEE Top Picks Honorable Mention

RESEARCH AND WORK EXPERIENCE

Intel, USA.

- **Graphics Research Organization (GRO)** *May 2023 - Present*
Graduate Research Intern (Full-time)
Managers: Tobias Zirr, Anton Sochenov, Anton Kaplanyan.
- **Next-Gen Architecture (NGA)** *August 2022 - May 2023*
Graduate Research Intern (Part-time)
Manager: Maxim Kazakov.
- **Graphics Research Organization (GRO)** *June 2022 - August 2022*
Graduate Research Intern (Full-time)
Managers: Rama Harihara, Anton Kaplanyan.
- **Cloud System Research Lab (CSR)** *Dec 2021 - May 2022*
Graduate Research Intern (Part-time)
Managers: Nilesh Jain, Ravishankar Iyer.
- **Heterogeneous Platforms Lab (HPL)** *May 2021 - Aug 2021*
Graduate Research Intern (Full-time)
Manager: Tanay Karnik.

Coordinated Science Laboratory (CSL), UIUC, USA *Jan 2019 - Present*
Research Assistant
Advisor: Prof. Rakesh Kumar

- Space Microdatacenters (**MICRO-56th, 2023**)
 - The paper identified a data transmission problem due to the increased volume of data generated by Earth observation (EO) satellites, a problem that cannot be solved by traditional data reduction techniques or RF-based infrastructure enhancement.
 - The proposed solution is to relocate the computation that would have happened on Earth into space, reducing the need for data transfer to Earth.
 - Analysis reveals small satellites currently in use can't meet the computation and power requirements of EO applications.
 - Advocated for space microdatacenters - large computational satellites designed to support in-space computation of EO data, with one 4KW space microdatacenter being able to support the computation needs of most applications, particularly when used with early discard.
 - Proposed co-design strategies to address communication bottleneck between EO satellites and space microdatacenters.
- Hardware Acceleration of Stable Diffusion (Ongoing)
- ML accelerators for accelerating gaming and vision applications (Ongoing)
- Hardware Acceleration of Neural Graphics (**ISCA-50th, 2023**)
 - Neural graphics as a potent replacement for traditional rendering algorithms.
 - In this paper we addressed the question does NG need HW support?
 - Disparity between desired performance and the current state of the art (1.5x-55x).
 - Identification of input encoding and MLP kernels as performance bottlenecks.
 - Proposal of NGPC a hardware solution to accelerate input encoding and MLP.
 - The NGPC achieves up to 58X end-to-end application-level performance improvement.
 - NGPC enables rendering of 4k Ultra HD resolution frames at 30 FPS for NeRF and 8k Ultra HD frames at 120 FPS for all other neural graphics applications.
- Understanding Interactions Between Chip Architecture and Uncertainties in Semiconductor Supply and Demand (Under Submission - Preprint available on arXiv.)
- Exploiting Short Application Lifetimes for Low Cost Hardware Encryption in Flexible Electronics (**DATE 2023**)
- Rethinking Programmable Earable Processors (**ISCA-49th, 2022**)
 - We first explore the hardware requirements for earable computing platforms like earphones, hearing aids, and smart glasses.
 - we propose EarBench, a suite of earable applications, to analyze performance gaps between earable computational needs and the capabilities of existing microprocessors.
 - Analysis reveals a performance gap of 13.54x-3.97x on average, with more complex microprocessors being energy-inefficient for earable applications.
 - EarBench applications are found to be dominated by a few DSP and ML-based kernels with significant computational similarity.
 - The proposed solution is SpEaC, a reconfigurable spatial architecture optimized for energy-efficient execution of earable kernels at the cost of generality.

- SpEaC outperforms modeled programmable cores (M4, M7, A53, HiFi4 DSP) by up to 99.3x, and offers substantial energy efficiency benefits, outperforming a low power Mali T628 MP6 GPU by 15.7x - 1087x.
- Model-specific Design of Deeply-Embedded Tiny Neural Network Accelerators (Ongoing).
 - Current edge NN accelerator hardware designs tend to be model-agnostic and are, therefore, over designed and expensive for a specific model.
 - For a suite of tiny neural networks, we show that customizing hardware to model structure can provide $101\times$ benefits in terms of inference throughput and $69\times$ benefits in terms of energy efficiency over Eyeriss. Additional $4.3\times$ and $2.4\times$ benefits are achievable in terms of area and energy efficiency respectively if model data is also known during design time.
- Printed Machine Learning Classifiers (**MICRO-53rd, 2020**)
 - Printed electronics meet cost, conformity, and non-toxicity needs unfulfilled by silicon-based systems in many applications.
 - We first explore the accuracy and cost of classification algorithms for two printed technologies, EGT and CNT-TFT.
 - Our EGT-bespoke Decision Tree and SVM classifiers exhibit 4x, 75x, 48x and 1.4x, 12.7x, 12.8x improvements in delay, power, and area, respectively, over conventional systems.
 - Our lookup-based classifiers enhance area and power efficiency by 1.93x and 1.65x with a 50% delay overhead for EGT Decision Trees. SVMs present a 40% delay cut and 8% and 1% area and power improvements.
 - EGT analog Decision Trees and SVMs surpass their digital equivalents in area and power by 437x, 27x and 490x, 1212x, respectively, but are 1.63x slower.
 - We presented several functional printed prototypes.
- Printed Microprocessors (**ISCA-47th, 2020.**)
 - We first explore the design space for microprocessors implemented in printing technologies, also devising standard cell libraries for these technologies.
 - Our design space exploration of printed microprocessor architectures across various parameters reveals our best cores surpass pre-existing ones by at least an order of magnitude in power and area.
 - Introducing printed-specific optimizations and Program-specific ISA (TP-ISA), we achieve up to 4.18x and 1.93x improvements in power and area, respectively. A Cross point-based instruction ROM outperforms a RAM-based design by 5.77x, 16.8x, and 2.42x in power, area, and delay, respectively.

National Electronics Complex of Pakistan (NECOP), Pakistan

Jul 2017 - Oct 2018

RF and Microwave Design Engineer - Assistant Manager

- Designed and Developed IMF Receiver, Antennas for Phased Array Radar, Branched Line Hybrid Coupler, Power Dividers and Band Pass Filters for X-Band Applications.

Research Internee at MERL - NUST, Pakistan

June 2016- July2017

- Designed and Developed a "Low Cost Short Range Frequency Modulated Continuous Wave Radar (FMCW) Radar". (**Second Position in Rector's Gold Medal Competition - NUST.**)

PUBLICATIONS

- 1 N. Bleier **M. Mubarik**, G. Swanson, R.Kumar "Space Microdatacenters" **MICRO-56th, 2023.**
- 2 **M. Mubarik**, R. Kanungo, T. Zirr, R.Kumar "Hardware Acceleration of Neural Graphics" **ISCA-50th, 2023.**
- 3 N. Bleier, **M. Mubarik**, S. Balaji, F. Rodriguez, A. Sou, S. White and R. Kumar, "Exploiting Short Application Lifetimes for Low Cost Hardware Encryption in Flexible Electronics," **DATE 2023.**
- 4 N. Bleier, **M. Mubarik**, S. Chakraborty, S. Kishore, R. Kumar, "Rethinking Programmable Earable Processors," **ISCA-49th, 2022.**
- 5 **M. Mubarik**, D. Weller, N. Bleier, M. Tomei, J. Aghassi-Hagmann, M. Tahoori, R. Kumar, "Printed Machine Learning Classifiers," **MICRO-53rd 2020, IEEE Micro Top Picks - Honorable Mention 2021.**
- 6 N. Bleier*, **M. Mubarik***, F. Rasheed*, J. Aghassi-Hagmann, M. Tahoori, R. Kumar, "Printed Microprocessors," **ISCA-47th, 2020.** [* Co-First Authors, Listed in Alphabetical Order] **Selected for the retrospective of the years 1996 through 2020 on the 50th anniversary of ISCA.**
- 7 R. Kanungo, S. Siva, N. Bleier, **M. Mubarik**, L. Varshney, R. Kumar "Understanding Interactions Between Chip Architecture and Uncertainties in Semiconductor Supply and Demand" - Under Submission, preprint arXiv:2305.11059, 2023.

TALKS

- Presented "Hardware Acceleration of Neural Graphics" paper at ISCA-50th, 2023.
- Gave a talk on "Neural Graphics: An Architectures Perspective" at ECE498SJP: Accelerator Architectures class at UIUC [2023].
- Gave a talk titled "Hardware Acceleration Opportunities in Neural Radiance Fields" at Intel [2022].
- Presented "Printed Machine Learning Classifiers" paper at MICRO-53rd, 2020.
- Gave a talk on "Hardware for Deep Learning" at ECE511: Advanced Computer Architecture class at UIUC [2021].

SELECT ACADEMIC ACHIEVEMENTS AND AWARDS

🏆 ISCA-50th 25 year retrospective (1996-2020), 2023.

- Selected for the retrospective of the years 1996 through 2020 on the 50th anniversary of ISCA for our paper "Printed Microprocessors" (ISCA-47th, 2020).

🏆 Machine Learning and Systems Rising Star, 2023.

- Harvard, Google, Nvidia and ML Commons.

🏆 IEEE Micro Top Picks - Honorable Mention, 2021.

- Our paper "Printed Machine Learning Classifiers" (MICRO-53rd, 2020) ranks among the Top 24 papers published in computer architecture venues in 2020.

🏆 Prime Minister's Silver Medal, NUST-CEME, 2018.

- 2nd highest CGPA (3.95/4.00) in the Electrical Engineering Department at NUST-CEME.

TECHNICAL SKILLS

- Programming languages: C++, Python, Verilog, SystemVerilog, CUDA, GLSL, Scala, SystemC, Tcl, Bash, Makefile
- Software: PyTorch, TensorFlow, TensorRT, Onnx-Runtime, Caffe, Keras, Synopsys-VCS, MATLAB, LATEX, Vim
- Frameworks: CUDA, Nvidia Nsight (Compute/Systems/Graphics), Vulkan, OpenGL, CHISEL, gem5.

COURSES

- 📖 ECE498SJP: Accelerator Architectures (Sanjay Patel) - Spring 2023.
- 📖 CS534: Advanced Topics in Computer Architecture (Sarita Adve) - Spring 2022.
- 📖 ECE598JH: Advanced Memory and Storage Systems (Jian Huang) - Fall 2021.
- 📖 CS598DHK: 3D Vision (Derek Hoiem) - Fall 2021.
- 📖 CS498ME Architecture for Mobile and Edge Computing (Saugata Ghose) - Spring 2021.
- 📖 CS598LCE Languages and Compilers for Edge Computing (Vikram S. Adve) - Spring 2021.
- 📖 ECE598NSG: Deep Learning in Hardware (Naresh Shanbhag) - Fall 2020.
- 📖 ECE524/CS563: Advanced Computer Security (Adam Bates) - Spring 2020.
- 📖 CS598SVA: App-Cust Heterogeneous Systems (Sarita Adve) - Spring 2020.
- 📖 CS598JT: Energy Efficient Computer Architecture (Josep Torrellas) - Fall 2019.
- 📖 ECE511: Advanced Computer Architecture (Rakesh Kumar) - Spring 2019.
- 📖 CS533: Parallel Computer Architecture (Josep Torrellas) - Spring 2019.

MEMBERSHIPS

- Member, Computer Architecture Student Association (CASA) Jan 2021 - Present
- Reviewed papers for ISCA, MICRO, HPCA and ASPLOS.